

Programming Strategies for Parallel Haskell



Applicant : Andrea Senese
Supervisor : Ugo de' Liguoro

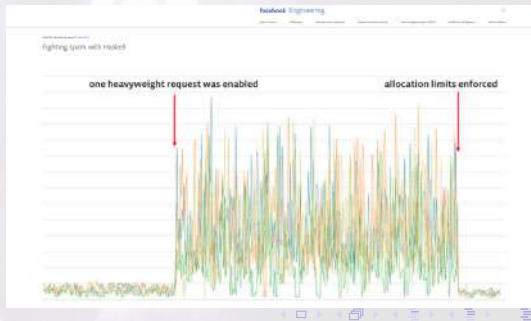
University of Turin
Mathematical, Physical and Natural Sciences School
Computer Science Department
Bachelor's Degree in Computer Science - Languages and Systems

- 1 *What is Haskell?*
- 2 *Determinism in Haskell*
- 3 *Monad for Non-Deterministic Action*
- 4 *Parallel Haskell : Eval Monad and Strategies*
- 5 *Par Monad*
- 6 *Parallel Type Inference Algorithm*

What is Haskell?



- *Haskell* is a **purely functional advanced programming language** (based on the **lambda calculus system**) that takes its name from the US mathematical-logical Haskell Curry;
- Suitable for **teaching, research, applications** and the **construction of large systems**;
- Fully described through the publication of a **syntax** and **formal semantics**;
- Haskell uses a lazy evaluation and it is a lazy language;
- So, without having **side effects** in **Haskell**, we can say that it is **deterministic** and therefore by its nature we can create **parallel** programs.



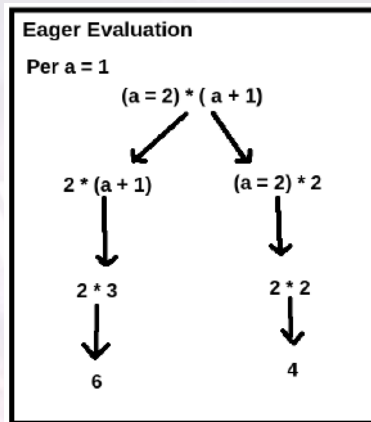
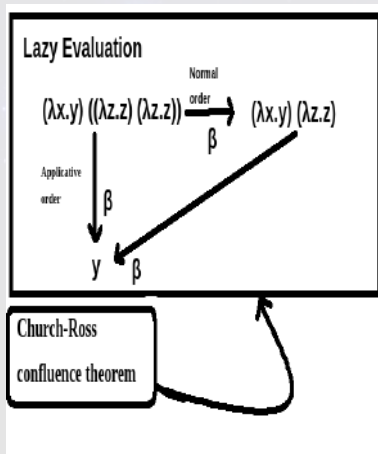


Figure 1: *Lazy Evaluation vs Eager Evaluation*

- Given a term M that can be reduced in one or more computation steps in N_1 and N_2 , then there exists N such that N_1 is reduced to N (in one or more computation steps) and N_2 is reduced to N (in one or more steps). Formally:

Theorem (Church-Ross Confluence)

If $M \rightarrow N_1$ and $M \rightarrow N_2$ then $\exists N$ such that $N_1 \rightarrow N$ and $N_2 \rightarrow N$

Indirect Proof:

Supposed that N_1 and N_2 are two normal forms of M then by the confluence theorem since we know that N_1 is reduced to N (in 0 steps) and also N_2 then for transitivity:

$$(N_1 = N \wedge N_2 = N) \Rightarrow N_1 = N_2$$

Monad for Non-Deterministic Action(side effect)

- *Monads in Haskell* can be thought of as composable computation descriptions;
- The essence of *Monad* is thus separation of composition timeline from the composed computation's execution timeline, as well as the ability of computation to implicitly carry extra data, as pertaining to the computation itself, in addition to its one (hence the name) output, that it will produce when run (or queried, or called upon). This allows monads supplementing pure calculations with features like *I/O*, *common environment*, *updatable state*, etc..;
- *Monads* are composed by three features:
 - ▶ a *type constructor* *M*;
 - ▶ A *bind* operator (*>>=*);
 - ▶ a *return* operator.

Monad

```
class Monad m where
    (>>=) :: m a -> (a -> m b) -> m b
    (>>)  :: m a -> m b                -> m b
    return :: a                        -> m a
    fail   :: String -> m a
```




Figure 2: *Monads*



Figure 3: *Parallelism* \neq *Concurrency*

- *Parallelism* and *Concurrency* are distinct concepts:
 - ▶ *Parallelism* means that several threads work on the same task to reduce computation calculus and throughput time \Rightarrow *Determinism*;
 - ▶ *Concurrency* consists of performing different tasks from different threads and sharing the CPU resource \Rightarrow *Non-Determinism*.
- *Haskell* offers the following tools to the programmer to parallelize their programs:
 - ▶ *Eval Monad* and *Strategies*;
 - ▶ *Par Monad*.

- *Eval Monad* is the first approach to express *Parallelism* in *Haskell* from *Control.Parallel.Strategies*, refine use of the function *rpar* and *rseq* functions. There are the basic operator in *Eval Monad*:
 - ▶ *runEval* → compute the operation in the *Eval Monad* and *return* the result;
 - ▶ *rpar* → the argument could be evaluated in parallel;
 - ▶ *rseq* → forces the sequential evaluation and waits for the result.
- For example :




Figure 4: *rpar/rpar* and *rpar/rseq*

- The solution is:

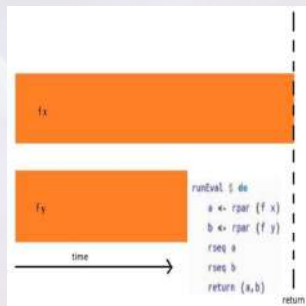


Figure 5: *rpar/rpar/rseq/rseq*

- *Eval Monad* can be used to diagnose programs with *Threadscope*.

- **Threadscope** was originally created to help programmers visualize the parallel execution of **Haskell** programs compiled with **GHC**. The idea is that when running a parallel **Haskell** program, the **Haskell** runtime system writes date and time reports of significant events to a log file. The **Threadscope** tool then reads this log file, and graphically shows the user what each **CPU** was doing in that time interval. The diagrams it shows reveal to the programmer where the program is using parallelism, and where it is not.

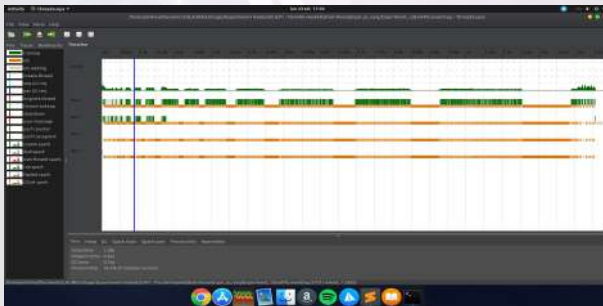


Figure 6: Threadscope tool

- There's a close relationship between the *Eval Monad* and *Strategies* as we can see in the *Haddock* documentation. *Strategies* step in as a higherlevel abstraction to ease many of the problems associated with regulating granularity and forcing evaluation manually;
- *Evaluation Strategies*, or simply *Strategies*, are ways used to modularize parallel code by separating the algorithm from the parallelism and provide ways to express parallel computations. *Strategies* have the following key features:
 - ▶ *Strategies* express deterministic parallelism: the result of the program is not affected by evaluating in parallel. The parallel tasks evaluated by a *Strategy* may have no side effects;
 - ▶ *Strategies* allow to separate the description of the parallelism from the logic of the program, enabling modular parallelism. The basic idea is to build a lazy data structure representing the computation, and then write a *Strategy* that describes how to traverse the data structure and evaluate components of it sequentially or in parallel;
 - ▶ *Strategies* are compositional: larger strategies can be built by gluing together smaller ones.

- For example : for evaluating list in parallel

ParList

```
parMap :: (a -> b) -> [a] -> [b]
parMap f xs = using (map f xs) (parList rseq)
```

```
evalList :: Strategy a -> Strategy [a]
evalList strat [] = return []
evalList strat (x:xs) = do
  y <- strat x
  ys <- evalList strat xs
  return (y : ys)
```

```
parList :: Strategy a -> Strategy [a]
parList strat = evalList (rparWith strat)
```

- The Tools mentioned so far lead to high failure rates because *rpar* requires the programmer to understand the operational properties of the program that reside, in most cases, in the implementation;
- Another practice used to correct the main weaknesses is the *Monad Eval*, which unfortunately is still limited (and does not go very far);
- The *Monad Par* offers three basic operations:
 - ▶ *runPar* → compute the operation in the *Par Monad* and *return* the result;
 - ▶ *fork* → create new *Thread*;
 - ▶ *new* → create new *IVar*;
 - ▶ *get* → get the result written in the *IVar*;
 - ▶ *put* → writes the result in the *IVar*.




Figure 7: Par Monad Semantics

Warning

*The **IVar** to ensure determinism can only be written once.*

We assume that $\text{runPar } M \Downarrow N$ and $\text{runPar } M \Downarrow N'$, then $N = N'$, and furthermore that if $\text{runPar } M \Downarrow N$ then there is no sequence of transitions starting from $\text{runPar } M$ that reaches a state in which no reduction is possible. Informally, $\text{runPar } M$ should either produce the same value consistently, or produce no value consistently (which is semantically equivalent to \perp). First, observe that the system contains no transitions that take a full **IVar** to an empty **IVar**, or that change the contents of a full **IVar**. A non-deterministic result can only arise due to a race condition: two different orderings of reductions that lead to a different result. To observe non-determinism, there must be multiple reduction sequences leading to applications of the rule (**putEmpty**) for the same **IVar** with different values. There is nothing in the semantics that prevents this from happening, but our determinism argument rests on the (**runPar**) rule, which requires that the state at the end of the reduction consists of only **done** M for some M . That is, the rest of the state has completed or deadlocked and been garbage collected by rules (**GCReturn**), (**GCEmpty**), (**GCFull**), and (**GCDeadlock**). In the case where there is a choice between multiple puts, one of them will not be able to complete and will remain in the state, and thus the **runPar** transition will never be applicable.

If some subterm evaluates to \perp during execution of `runPar`, then the value of the whole `runPar` should be \perp . The sequential scheduler implements this, but the parallel scheduler given above does not. In practice, however, it will make little difference: if one of the workers encounters a \perp , then the likely result is a deadlock, because the thread that the worker was executing will probably have outstanding puts. Since deadlock is semantically equivalent to \perp , the result in this case is equivalent. Nevertheless, modifying the implementation to respect the semantics is not difficult, and we plan to modify our real implementation to do this in due course.

- For example : To parallelize a quicksort

ParQuick

```
parQuickSort :: (Ord a , NFData a) => [a] -> Par [a]
parQuickSort [] = return []
parQuickSort (x : xs) = do
    p1 <- spawn(parQuickSort(filter(< x) xs))
    p2 <- spawn(parQuickSort(filter(>= x) xs))
    left <- get p1
    right <- get p2
    return $ left ++ (x : right)

parQuickSortLimit :: (Ord a , NFData a) => Int -> [a] -> Par [a]
parQuickSortLimit 0 xs = return $ quickSortSeq xs
parQuickSortLimit d [] = return []
parQuickSortLimit d (x : xs) = do
    p1 <- spawn(parQuickSortLimit(d - 1)(filter(< x) xs))
    p2 <- spawn(parQuickSortLimit(d - 1)(filter(>= x) xs))
    left <- get p1
    right <- get p2
    return $ left ++ (x : right)
```

- This graph illustrates the speedup of the parallel quicksort when applying varying thresholds. The results were achieved on a quad-core CPU running 8 threads:




Figure 8: *Timeline ParQuickSortLimit*

- In *Threlescope*:




Figure 9: *Sequential quicksort as seen in Threlescope on quad-core CPU with 8-threads.*




Figure 10: *Parallel quicksort as seen in Threlescope on quad-core CPU with 8-threads.*

Warning

The input data is not necessarily split evenly in each recursion in some runs there can be bad pivot elements which infers that the work is divided very unevenly between the cores. This will slow down the execution!

Parallel Type Inference Algorithm



- An example that naturally fits the dataflow model is program analysis, in which information is typically propagated from definition sites to usage sites in a program;
- Type inference gives rise to a dataflow graph; each binding is a node in the graph with inputs corresponding to the free variables of the binding, and a single output represents the derived type for the binding. For example, the following set of bindings can be represented by the dataflow graph:




Figure 11: *Dataflow graph of bindings*


We define the function *infer* as follows:

parInfer

```
--(1)
parInfer :: [(Var, Expr)] -> [(Var, Type)]
parInfer bindings = runPar $ do
    let binders = map fst bindings --(2)
    ivars <- replicateM (length binders) new --(3)
    let Env = Map.fromList (zip binders ivars) --(4)
    mapM_ (fork . infer env) bindings --(5)
    types <- mapM_ get ivars --(6)
    return (zip binders types) --(7)
```

Parallel Type Inference Algorithm

A concrete execution :



- (1) call `parInfer -> parInfer [(var1,exp1),(var2,exp1),...,(var_n,exp_n)]`
- (2) let `binders = map fst [(var1,exp1),(var2,exp1),...,(var_n,exp_n)] = [var1,var2,...,var_n]` (forall elem in list apply `fst(elem)`)
- (3) `ivars = replicateM (length [(var1,exp1),(var2,exp1),...,(var_n,exp_n)]) new = replicateM n new = [Ivar1,Ivar2,..., Ivar_n]`
(create a list of length `n` with `n` `Ivar`)
- (4) `env = Map.fromList (zip [var1,var2,...,var_n] [Ivar1,Ivar2,...,Ivar_n]) = Map.fromList [(var1,Ivar1),(var2,Ivar2),...,(var_n,Ivar_n)] =`

Key	Value
1	var1
...	...
n	var_n

- (5) Infer calls put for each binding in the list for write the type in the `Ivar` mapped by that variable in env.
- (6) `type = mapM_get [Ivar1,...,Ivar_n] = [type1,type2,...,type_n]`
(wait the result computed by the childs).
- (7) return `(zip [var1,var2,...,var_n] [type1,type2,...,type_n]) = [(var1,type1),...,(var_n,type_n)]` (return the result)

- *This implementation extracts the maximum parallelism inherent in the dependency structure of the given set of bindings, with very little effort on the part of the programmer;*
- *The same trick can be pulled using lazy evaluation of course, and indeed we could achieve the same parallel structure using lazy evaluation together with Strategies, but the advantage of the Par monad version is that the structure is programmed explicitly and the runtime behaviour is not tied to one particular evaluation strategy (or compiler).*

- In *Threасcope*:



Figure 12: *ParInfer* as seen in *Threасcope* on quad-core CPU with 8-threads.



Figure 13: *ParInfer* as seen in *Threасcope* on quad-core CPU limiting it to 2 cores

GRAZIE A TUTTI PER L'ATTENZIONE

